🎯 Retail Intelligence Platform - Final Submission Guide

Team BNZ - Snowflake x Accenture Hackathon

✅ Complete Implementation Checklist

Phase 1: Environment Setup ☐

- [ ] Snowflake trial account created (US West - Oregon)
- [ ] Database `RETAIL_INTELLIGENCE_DB` created
- [ ] Schemas created: RAW, STAGING, ANALYTICS, MART
- [ ] Warehouse `RETAIL_WH` created and configured
- [ ] ACCOUNTADMIN role used for permissions

Phase 2: Data Ingestion ☐

- [ ] Downloaded Abt-Buy dataset from https://dbs.uni-leipzig.de/files/datasets/Abt-Buy.zip
- [ ] Created raw tables: ABT_PRODUCTS, BUY_PRODUCTS, PRODUCT_MAPPING_TRUTH
- [ ] Created file format and stage
- [ ] Loaded all CSV files successfully
- [ ] Verified data counts match expected values
- [ ] Data quality checks passed

Phase 3: Data Transformation ☐

- [ ] DBT project setup (or staging views created)
- [ ] STG_ABT_PRODUCTS created with cleaned data
- [ ] STG_BUY_PRODUCTS created with cleaned data
- [ ] Price conversions working correctly
- [ ] Text normalization applied (UPPER, TRIM)
- [ ] Null handling implemented

Phase 4: Feature Engineering with Snowpark ☐

- [ ] Snowpark Python environment configured
- [ ] Custom UDFs created (Jaccard similarity, Levenshtein distance)
- [ ] Product pair features table created
- [ ] Features include: name similarity, description similarity, price differences
- [ ] Token overlap calculated
- [ ] Feature validation completed

Phase 5: AI/ML with Cortex ☐

- [ ] **PRODUCT_EMBEDDINGS** table created using EMBED_TEXT_768
- [ ] Embeddings generated for all products (ABT + BUY)
- [ ] **PRODUCT_MATCHES_AI** table created with cosine similarity
- [ ] Match confidence levels assigned (HIGH, MEDIUM, LOW)
- [ ] Match ranking implemented (top N per product)
- [ ] **PRODUCT_CATEGORIES** table created using COMPLETE LLM
- [ ] AI categorization applied to all products
- [ ] Semantic similarity threshold tuned (≥ 0.55)

Phase 6: Price Intelligence ☐

- [ ] **PRICE_INTELLIGENCE** view created
- [ ] Price differences calculated (absolute & percentage)
- [ ] Price positions identified (cheaper/parity)
- [ ] Competitiveness scores computed
- [ ] Category-level price analysis available
- [ ] Price opportunity identification working

Phase 7: Model Validation ☐

- [ ] **MATCH_VALIDATION** view created
- [ ] Ground truth comparison implemented
- [ ] Precision metrics calculated by confidence level
- [ ] **MODEL_PERFORMANCE** view created
- [ ] **OVERALL_PERFORMANCE** metrics available
- [ ] Accuracy > 70% for HIGH confidence matches
- [ ] Coverage metrics tracked

Phase 8: Cortex Analyst ☐

- [ ] Semantic model YAML file created
- [ ] All tables properly defined in YAML
- [ ] Dimensions and measures configured
- [ ] Relationships defined
- [ ] Sample questions included
- [ ] YAML uploaded to CORTEX_ANALYST_STAGE
- [ ] Cortex Analyst tested with natural language queries
- [ ] At least 5 test questions validated

Phase 9: Snowflake Intelligence Dashboards ☐

- [ ] Dashboard created in Snowsight
- [ ] Key metrics tiles added (match statistics, prices)
- [ ] Visualizations created (charts, graphs)

- [ ] Match confidence distribution chart
- [ ] Price position breakdown chart
- [ ] Category analysis charts
- [ ] Model performance metrics displayed
- [ ] Dashboard filters configured
- [ ] Screenshots captured

Phase 10: Streamlit Application ☐

- [ ] Streamlit app created in Snowflake
- [ ] Connection to Snowflake configured
- [ ] **Product Matching** tab implemented
- [ ] **Price Intelligence** tab implemented
- [ ] **Category Analysis** tab implemented
- [ ] **Search Products** tab implemented
- [ ] **Model Performance** tab implemented
- [ ] **AI Assistant** tab with Cortex Analyst
- [ ] Interactive filters working
- [ ] Data export functionality added
- [ ] Error handling implemented
- [ ] App tested end-to-end

📊 Snowflake Features Used - Verification

Required Features ✅

Feature	Used	Evidence
Snowpark	✅	Custom Python UDFs for similarity calculations
DBT Projects	✅	Staging and analytics models for data transformation
Cortex AI - Embeddings	✅	EMBED_TEXT_768 for semantic product vectors
Cortex AI - Complete	✅	LLM for product categorization
Cortex Agents	✅	Cortex Analyst for natural language queries
AI/SQL	✅	VECTOR_COSINE_SIMILARITY for product matching
Snowflake Intelligence	✅	Interactive dashboards in Snowsight
Streamlit	✅	Full-featured web application

Additional Features Used 🌟

- **Vector Search**: Semantic similarity using embeddings
- **Dynamic Tables**: (Optional) For real-time updates
- **Stages**: Internal stage for data loading
- **Views**: Analytics and mart layer views
- **UDFs**: Custom text processing functions
- **Time Travel**: Data versioning capability

📁 Deliverables Checklist

1. Code Repository ☐

```
retail-intelligence-platform/
├── README.md (comprehensive with setup instructions)
├── architecture_diagram.png
├── sql/
│   ├── 01_setup.sql
│   ├── 02_data_loading.sql
│   ├── 03_transformations.sql
│   ├── 04_ml_features.sql
│   ├── 05_cortex_ai.sql
│   └── 06_analytics.sql
├── snowpark/
│   ├── feature_engineering.py
│   ├── requirements.txt
│   └── test_features.py
├── dbt/ (if used)
│   ├── dbt_project.yml
│   └── models/
├── cortex/
│   ├── retail_semantic_model.yaml
│   └── test_analyst.py
├── streamlit/
│   ├── app.py
│   └── requirements.txt
├── dashboards/
│   ├── dashboard_queries.sql
│   └── screenshots/
├── tests/
│   ├── data_quality.sql
```

```
│   └── validation.sql
└── docs/
    ├── setup_guide.md
    ├── user_guide.md
    └── technical_documentation.md
```

2. Documentation ☐

☐ **README.md** with:

- Problem statement
- Solution overview
- Architecture diagram
- Setup instructions
- Features used
- Results and metrics
- Team information

☐ **Setup Guide** with:

- Step-by-step installation
- Prerequisites
- Configuration details
- Troubleshooting section

☐ **Technical Documentation** with:

- Data model diagrams
- SQL query explanations
- ML model details
- API documentation

☐ **User Guide** with:

- Dashboard navigation
- Streamlit app usage
- Cortex Analyst queries
- Feature descriptions

3. Presentation Materials ☐

☐ **PowerPoint/PDF Slides** (10-15 slides):

- Title slide with team name
- Problem statement

- Solution architecture
- Technical implementation
- Snowflake features showcase
- Demo screenshots
- Results and metrics
- Business value
- Future enhancements
- Q&A slide

☐ **Demo Video** (3-5 minutes):

- Introduction (30 sec)
- Problem & solution overview (1 min)
- Live demo of key features (2-3 min)
- Results and impact (30 sec)
- Conclusion (30 sec)
- Professional quality audio/video
- Screen recordings of actual application

4. Screenshots & Visuals ☐

☐ Architecture diagram
☐ Data flow diagram
☐ Snowsight dashboard screenshots (at least 3)
☐ Streamlit app screenshots (all tabs)
☐ Cortex Analyst query examples
☐ SQL query examples with results
☐ Model performance charts
☐ Price intelligence visualizations

5. Test Results ☐

☐ Data quality test results
☐ Model performance metrics:

- Overall accuracy: ___%
- Precision by confidence level
- Recall metrics
- Coverage statistics

☐ Performance benchmarks:

- Query response times

- Embedding generation time
- Dashboard load time

☐ Sample outputs:

- Top 10 product matches
- Price analysis results
- Category distribution
- AI-generated insights

🎯 Key Metrics to Highlight

Quantitative Results

1. **Data Scale**
   - Total products processed: ___
   - Total matches generated: ___
   - High confidence matches: ___
2. **Model Performance**
   - Overall precision: ___%
   - HIGH confidence precision: ___%
   - MEDIUM confidence precision: ___%
   - Recall rate: ___%
3. **Business Insights**
   - Average price difference: ___%
   - Products with >10% savings: ___
   - Categories analyzed: ___
   - Price opportunities identified: ___
4. **Technical Performance**
   - Average query time: ___ ms
   - Embedding generation time: ___ sec
   - Dashboard load time: ___ sec
   - Concurrent users supported: ___

Qualitative Results

1. **Innovation**
    - Multi-strategy matching approach
    - AI-powered categorization
    - Natural language query interface
    - Real-time price intelligence
2. **User Experience**
    - Intuitive dashboard design
    - Interactive filtering
    - Natural language queries
    - Export capabilities
3. **Business Value**
    - Automated product matching
    - Competitive pricing insights
    - Market intelligence
    - Decision support

🚀 Submission Guidelines

Before Submitting

1. **Final Testing**
    - Run all SQL scripts from scratch
    - Test Streamlit app end-to-end
    - Verify all dashboards load correctly
    - Test Cortex Analyst with 10+ queries
    - Check all exports and downloads
2. **Code Quality**
    - Add comments to complex queries
    - Include error handling
    - Use consistent formatting
    - Remove hardcoded credentials

- Add logging where appropriate

3. **Documentation Review**
    - Check for typos and grammar
    - Verify all links work
    - Ensure screenshots are clear
    - Update version numbers
    - Add contact information
4. **Performance Optimization**
    - Add appropriate indexes
    - Optimize slow queries
    - Cache frequently used results
    - Review warehouse sizing

Submission Checklist

- [ ] All code pushed to GitHub/GitLab
- [ ] README.md is complete and clear
- [ ] Documentation is comprehensive
- [ ] Presentation slides finalized
- [ ] Demo video recorded and uploaded
- [ ] Screenshots captured and organized
- [ ] Test results documented
- [ ] Requirements.txt/dependencies listed
- [ ] No sensitive data or credentials in code
- [ ] License file included (if required)

What to Submit

1. **GitHub Repository URL** with:
    - Public access enabled
    - Clear README
    - All code and documentation
2. **Presentation Deck** (PDF format)
3. **Demo Video** (MP4, hosted on YouTube/Vimeo or included in repo)
4. **Snowflake Account Details** (for judges to test):
    - Account URL
    - Read-only username/password

- Database and schema names
- Dashboard URLs
- Streamlit app URL

5. **Submission Form** with:
    - Team name: BNZ
    - Use case: Retail Intelligence Platform
    - Features used: (list all)
    - GitHub URL
    - Video URL
    - Any special instructions

🌟 Tips for Standing Out

Technical Excellence

1. **Advanced Features**
    - Implement multiple matching strategies
    - Use ensemble methods for categorization
    - Add confidence calibration
    - Implement A/B testing framework
2. **Performance Optimization**
    - Optimize vector similarity calculations
    - Use materialized views strategically
    - Implement caching layers
    - Add query profiling
3. **Code Quality**
    - Write clean, documented code
    - Use design patterns
    - Implement comprehensive testing
    - Add CI/CD pipeline

User Experience

1. **Dashboard Design**
    - Professional visualizations
    - Intuitive navigation
    - Responsive layout
    - Helpful tooltips
2. **Streamlit App**
    - Modern UI/UX
    - Fast response times
    - Error messages that guide users
    - Export functionality
3. **Documentation**
    - Clear setup instructions
    - Visual diagrams
    - Use case examples
    - Troubleshooting guide

Business Impact

1. **Quantify Value**
    - Cost savings identified
    - Time saved vs manual matching
    - Accuracy improvement
    - Scale achieved
2. **Real-World Application**
    - Explain use cases
    - Show scalability
    - Discuss deployment strategy
    - Address data privacy
3. **Future Roadmap**
    - Enhancement ideas
    - Additional features
    - Integration possibilities

- Scaling strategy

📞 Support Resources

Snowflake Documentation

- Cortex AI: https://docs.snowflake.com/en/user-guide/snowflake-cortex
- Snowpark: https://docs.snowflake.com/en/developer-guide/snowpark/python/index
- Streamlit: https://docs.streamlit.io/
- DBT: https://docs.getdbt.com/

Community Support

- Snowflake Community: https://community.snowflake.com/
- Stack Overflow: Tag [snowflake-cloud-data-platform]
- GitHub Discussions

Hackathon Specific

- Review hackathon rules and judging criteria
- Check submission deadline
- Verify required features are used
- Confirm submission format

🏆 Final Checklist Before Submission

- [ ] All features from requirements are implemented
- [ ] Code runs without errors
- [ ] Documentation is complete
- [ ] Presentation is polished
- [ ] Demo video is ready
- [ ] Test results are documented
- [ ] Repository is public
- [ ] All team members reviewed
- [ ] Submission form filled out
- [ ] Deadline confirmed

- [ ] Backup copy saved

Good Luck! 🎉

You've built an impressive AI-powered Retail Intelligence Platform using Snowflake's cutting-edge features. Make sure to:

1. **Test everything** one final time
2. **Document thoroughly** - judges love clear docs
3. **Show business value** - not just technical prowess
4. **Practice your demo** - first impressions matter
5. **Submit early** - don't wait until the last minute

Remember: Your solution solves a real business problem using advanced AI and modern data engineering practices. Be confident and highlight the innovation!

Questions or Issues?

If you encounter any problems:

1. Check the Snowflake documentation
2. Review error messages carefully
3. Test queries in isolation
4. Use Snowflake's query profiler
5. Check community forums

Team BNZ - You've got this! 🚀